

Chapter 4. (Part 2)

The Greedy Approach

경북대학교 배준현 교수

[\(joonion@knu.ac.kr\)](mailto:joonion@knu.ac.kr)



4.1 Minimum Spanning Trees

- Prim's Algorithm
- Kruskal's Algorithm

4.2 Dijkstra's Algorithm for Single-Source Shortest Paths

4.3 Scheduling Problem

4.4 Huffman Code

4.5 The Knapsack Problem: Greedy .vs. Dynamic Programming



4.3 Scheduling Problem

- The *Scheduling* Problem:
 - The *time in the system* is
 - the time spent both waiting and being served.
 - The problem of *minimizing* the *total time in the system*
 - has many applications.
 - ex) scheduling users' access to a bank counter or a disk drive.
 - You would learn it in detail
 - when you study the schedulers of operating system.



4.3 Scheduling Problem

■ *Scheduling with Deadlines:*

- Another scheduling problem
 - occurs when *each job* takes the *same amount of time* to complete,
 - but has a *deadline* by which it must start to *yield a profit*
 - associated with the job.
- The goal is
 - to schedule the jobs to *maximize* the *total profit*.



4.3 Scheduling Problem

- The Problem of *Scheduling with Deadlines*:
 - Each job *takes one unit* of time to finish
 - and has a *deadline* and a *profit*.
 - If the job starts *before or at* its deadline, the profit is obtained.
 - Therefore, not all jobs have to be scheduled.
 - A schedule is called ***impossible***,
 - if a job is scheduled *after its deadline*.
 - We need not consider any impossible schedule
 - because the job in that schedule does not yield any profit.



4.3 Scheduling Problem

Job	Deadline	Profit
1	2	30
2	1	35
3	2	25
4	1	40

1	2	3	4
---	---	---	---

Job 1 **Job 2** (impossible)

Job 1 Job 3

Job 1 **Job 4** (impossible)

Job 2 Job 1

Job 2 Job 3

Job 2 **Job 4** (impossible)

Schedule	Total Profit
[1, 3]	55
[2, 1]	65
[2, 3]	60
[3, 1]	55
[4, 1]	70 (optimal)
[4, 3]	65

$$\text{profit}([1, 3]) = 30 + 25 = 55$$

$$\text{profit}([4, 1]) = 40 + 30 = 70$$



4.3 Scheduling Problem

- The Greedy Approach to the Problem:
 - To consider all schedules, a brute-force approach,
 - takes *factorial* time. (worse than exponential)
 - A reasonable greedy approach for solving the problem would be:
 - First, *sort* the jobs in *non-increasing* order *by profit*.
 - Next, *inspect* each job in sequence
 - and *add* it to the *schedule* if it is *possible*.



4.3 Scheduling Problem

- The Greedy Approach to the Problem:
 - A *sequence* is called a **feasible sequence**
 - if all the jobs in the sequence *start by their deadlines*.
 - ex) [4, 1]: feasible sequence, [1, 4]: not a feasible sequence.
 - A *set of jobs* is called a **feasible set**
 - if there exists *at least one feasible sequence* for the jobs in the set.
 - ex) {1, 4}: feasible set, {2, 4} not a feasible set.
 - Our goal is to find an **optimal** sequence,
 - which is *a feasible sequence with maximum total profit*.
 - An **optimal set of jobs** is the set of jobs in an optimal sequence.



4.3 Scheduling Problem

- High-level greedy algorithm for the problem:

sort the jobs in nonincreasing order by profit;

$S = \emptyset$;

while (*the instance is not solved*) {

 select next job;

if (*S is feasible with this job added*)

 add this job to S ;

if (*there are no more jobs*)

 the instance is solved;

}



4.3 Scheduling Problem

Job	Deadline	Profit
1	3	40
2	1	35
3	1	30
4	3	25
5	1	20
6	3	15
7	2	10

1. $S = \phi$
2. $S = \{1\}$, [1] is feasible
3. $S = \{1, 2\}$, [2, 1] is feasible
4. $S = \{1, 2, 3\}$, rejected
there is no feasible sequence for this set : $S = \{1, 2\}$
5. $S = \{1, 2, 4\}$, [2, 1, 4] is feasible
6. $S = \{1, 2, 4, 5\}$, rejected
 $S = \{1, 2, 4\}$
7. $S = \{1, 2, 4, 6\}$, rejected
 $S = \{1, 2, 4\}$
8. $S = \{1, 2, 4, 7\}$, rejected
 $S = \{1, 2, 4\}$ (feasible set) [2, 1, 4] (feasible sequence)



4.3 Scheduling Problem

- An efficient way to *determine* whether a set is *feasible*:
 - Lemma:
 - Let S be a set of jobs, then S is *feasible*
 - if and only if the sequence obtained by ordering
 - the jobs in S according to *nondecreasing deadlines*
 - is *feasible*.

$S = \{1, 2, 4, 7\}$: feasible?

not feasible

We need only check the feasibility of the sequence:

[2, 7, 1, 4]



1



2



3



3

not feasible



4.3 Scheduling Problem

ALGORITHM 4.4: Scheduling with Deadlines

```
void schedule(int n, int deadline[], sequence_of_integer &J) {  
    int i;  
    sequence_of_integer K;  
  
    J = [1];  
    for (i = 2; i <= n; i++) {  
        K = J with i added according to nondecreasing values of deadline[i];  
        if (K is feasible)  
            J = K;  
    }  
}
```



4.3 Scheduling Problem

Job	Deadline	Profit
1	3	40
2	1	35
3	1	30
4	3	25
5	1	20
6	3	15
7	2	10

The jobs are already sorted by the profit

1. $J = [1]$
2. $K = [2,1]$, K is feasible
 $J = [2,1]$
3. $K = [2,3,1]$ is rejected, because K is not feasible
4. $K = [2,1,4]$, K is feasible
 $J = [2,1,4]$
5. $K = [2,5,1,4]$ is rejected
6. $K = [2,1,4,6]$ is rejected
7. $K = [2,7,1,4]$ is rejected

- $J = [2, 1, 4]$ is the final result
- $Total Profit = 35 + 40 + 25 = 100$



4.3 Scheduling Problem

```
typedef vector<int> sequence_of_integer;

bool is_feasible(sequence_of_integer& K, sequence_of_integer& deadline) {
    for (int i = 1; i < K.size(); i++)
        if (i > deadline[K[i]])
            return false;
    return true;
}
```



4.3 Scheduling Problem

```

void schedule(int n, sequence_of_integer& deadline, sequence_of_integer &J) {
    sequence_of_integer K;
    J.clear();
    J.push_back(0); // for an empty job
    J.push_back(1);
    for (int i = 2; i <= n; i++) {
        // K = J with i added according to nondecreasing values of deadline[i];
        K.resize(J.size());
        copy(J.begin(), J.end(), K.begin());
        int j = 1;
        while (j < K.size() && deadline[K[j]] <= deadline[i])
            j++;
        K.insert(K.begin() + j, i);
        if (is_feasible(K, deadline)) {
            // J = K
            J.resize(K.size());
            copy(K.begin(), K.end(), J.begin());
        }
    }
}

```



4.3 Scheduling Problem

- Time Complexity of Algorithm 4.4 (Worst-Case)
 - Basic Operation: the operation of comparisons to *sort*, to do $K = J$, and to *check* if K is *feasible*.
 - Input Size: n , the *number of jobs*.
 - In each iteration of the for- i loop, we need to do
 - at most $i - 1$ comparisons to add the i th job of K ,
 - and at most i comparisons to *check* if K is *feasible*.
 - Therefore, the worst case is

$$W(n) = \sum_{i=2}^n [(i-1) + i] = n^2 - 1 \in \Theta(n^2)$$



4.4 Huffman Code

- The problem of *data compression*
 - is to find an *efficient method* for *encoding a data file*.
 - A *binary code* is a common way to represent a data file.
 - A *codeword* is a *unique binary string*
 - representing *each character* in a binary code.
 - A *fixed-length* binary code
 - represents each character using the *same number of bits*.
 - A *variable-length* binary code is a more efficient coding.
 - It represents different characters using *different number of bits*.



4.4 Huffman Code

- The Problem of the *Optimal Binary Code*:
 - Given a file (or a string of characters),
 - find a *binary character code* for the characters in the file,
 - which represents the file in the *least number of bits*.
 - We discuss the encoding method, called *Huffman code*,
 - then we develop *Huffman's algorithm* for solving this problem.



4.4 Huffman Code

File = ababcbbbc, character set = {a, b, c}

Character	Fixed-length Binary Code
a	00
b	01
c	10

a b a b c b b b c
00 01 00 01 10 01 01 01 10

- It takes **18 bits** with this encoding

Character	Variable-length Binary Code
a	10
b	0
c	11

a b a b c b b b c
10 0 10 0 11 0 0 0 11

- 'b' occurs *most frequently*:
- Encode 'b' with one bit (0)
- Encode 'a' and 'c' starting with 1 to distinguish from 'bb'
- 'a': 10, 'c': 11
- It takes only **13 bits** with this encoding




4.4 Huffman Code

■ Prefix Code

- is one particular type of variable-length code.
- In a prefix code, *no codeword* for one character
 - constitutes the *beginning of the codeword* for another character.
- Every prefix code can be represented by
 - a *binary tree* whose *leaves* are *the characters* that are to be encoded.
- The advantage of a prefix code is that
 - we *need not look ahead* when parsing the file.

- b: 0
- a: 10
- c: 11






4.4 Huffman Code

$S = \{a, b, c, d, e, f\}$


Character	Frequency	$Code_1$	$Code_2$	$Code_3$
a	16	000	10	00
b	5	001	11110	1110
c	12	010	1110	110
d	17	011	110	01
e	10	100	11111	1111
f	25	101	0	10

- $Code_1$: Fixed-Length
- $Code_3$: Huffman code

$Code_2$



$Code_3$





4.4 Huffman Code

- Computing the *number of bits* for *encoding*:
 - Given the binary tree T corresponding to some code
 - the number of bits it takes to encode a file is given by
 - $bits(T) = \sum_{i=1}^n frequency(v_i) \times depth(v_i)$
 - where $\{v_1, v_2, \dots, v_n\}$ is the set of characters in the file,
 - $frequency(v_i)$ is the number of times v_i occurs in the file,
 - and $depth(v_i)$ is the depth of v_i in T .
- $bits(Code_1) = 255$
 - $bits(Code_2) = 231$
 - $bits(Code_3) = 212$



4.4 Huffman Code

■ Huffman's Algorithm

- *Huffman* developed a *greedy* algorithm
 - that produces an *optimal binary character code* by constructing
 - a *Huffman code*, a *binary tree* corresponding to an *optimal code*.
- We need a *type declaration* for the node of binary tree.
- We also need to use a *priority queue*
 - in which the character with the *lowest frequency* is *removed next*.
 - It can be implemented as a *min-heap*.



4.4 Huffman Code

- High-level pseudo-code for the Huffman's algorithm:

```
n = number of characters in the file;
Arrange n pointers to nodetype records in a PQ;

for (i = 1; i <= n - 1; i++) {
    remove(PQ, p);
    remove(PQ, q);
    r = new nodetype;
    r->left = p;
    r->right = q;
    r->frequency = p->frequency + q->frequency;
    insert(PQ, r);
}
remove(PQ, r);
return r;
```




4.4 Huffman Code

```
typedef struct node *node_ptr;
typedef struct node {
    char symbol;    // the value of a character.
    int frequency;  // the number of times the character is in the file.
    node_ptr left;
    node_ptr right;
} node_t;

struct compare {
    bool operator()(node_ptr p, node_ptr q) {
        return p->frequency > q->frequency;
    }
};

typedef priority_queue<node_ptr, vector<node_ptr>, compare> PriorityQueue;
```







4.4 Huffman Code

```
void huffman(int n, PriorityQueue& PQ)
{
    for (int i = 1; i <= n - 1; i++) {
        node_ptr p = PQ.top(); PQ.pop();
        node_ptr q = PQ.top(); PQ.pop();
        node_ptr r = create_node('+', p->frequency + q->frequency);
        r->left = p;
        r->right = q;
        PQ.push(r);
    }
}
```





4.4 Huffman Code





4.4 Huffman Code





4.4 Huffman Code


- Time Complexity of Huffman's Algorithm
 - If a priority queue is implemented *as a min-heap*,
 - each heap operation (insert & remove) requires $\Theta(\lg n)$ time.
 - Since there are $n - 1$ passes through the for- i loop,
 - the algorithm runs in $\Theta(n \lg n)$ time.

- It is *provable* that
 - Huffman's algorithm always produces an optimal binary code.
 - based on Lemma:
 - The binary tree corresponding to an *optimal binary prefix code* is *full*.
 - That is, *every nonleaf node has two children*.



4.4 Huffman Code

- How to *encode* a string of characters?
 - Make a *hash map* of characters and binary codes, then,
 - Given with a string of characters, *afdbce*,
 - you can *transform* it directly into *00100111101101111*.



a → 00
 d → 01
 f → 10
 c → 110
 b → 1110
 e → 1111

a f d b c e
 ↓ ↓ ↓ ↓ ↓ ↓
 00 10 01 1110 110 1111





4.4 Huffman Code

```
void make_encoder(node_ptr node, string code, map<char, string> &encoder) {  
    if (node->symbol != '+') { // leaf node  
        encoder[node->symbol] = code;  
    } else { // internal node  
        make_encoder(node->left, code + "0", encoder);  
        make_encoder(node->right, code + "1", encoder);  
    }  
}
```



4.4 Huffman Code

- How to *decode* an encoded binary string?
 - Given with an encoded binary string, **00100111101101111**,
 - you can *traverse* the binary tree to decode it into *afdbce*.





4.4 Huffman Code

```
void decode(node_ptr root, node_ptr node, string s, int i) {  
    if (i <= s.length()) {  
        if (node->symbol != '+') { // leaf node  
            cout << node->symbol;  
            decode(root, root, s, i);  
        } else { // internal node  
            if (s[i] == '0')  
                decode(root, node->left, s, i + 1);  
            else // s[i] == '1'  
                decode(root, node->right, s, i + 1);  
        }  
    }  
}
```

Any Questions?

